

Class07Lab

Joshua Mac

2025-02-03

PCA of UK food data

First, read the .csv to open our data on food in the UK.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  5
```

A1. So there are 17 rows and 5 columns.

To preview the first 6 of the 17 rows, we'll use:

```
## View all we can use View(x), which is case-sens btw
head(x)
```

| | X | England | Wales | Scotland | N.Ireland |
|---|---------------|---------|-------|----------|-----------|
| 1 | Cheese | 105 | 103 | 103 | 66 |
| 2 | Carcass_meat | 245 | 227 | 242 | 267 |
| 3 | Other_meat | 685 | 803 | 750 | 586 |
| 4 | Fish | 147 | 160 | 122 | 93 |
| 5 | Fats_and_oils | 193 | 235 | 184 | 209 |
| 6 | Sugars | 156 | 175 | 147 | 139 |

Our first column was taken as the row headers, which isn't ideal. Let's get rid of it. WE'll use `rownames()` to set that up.

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

| | England | Wales | Scotland | N.Ireland |
|---------------|---------|-------|----------|-----------|
| Cheese | 105 | 103 | 103 | 66 |
| Carcass_meat | 245 | 227 | 242 | 267 |
| Other_meat | 685 | 803 | 750 | 586 |
| Fish | 147 | 160 | 122 | 93 |
| Fats_and_oils | 193 | 235 | 184 | 209 |
| Sugars | 156 | 175 | 147 | 139 |

Now, `dim(x)` should be:

```
dim(x)
```

```
[1] 17  4
```

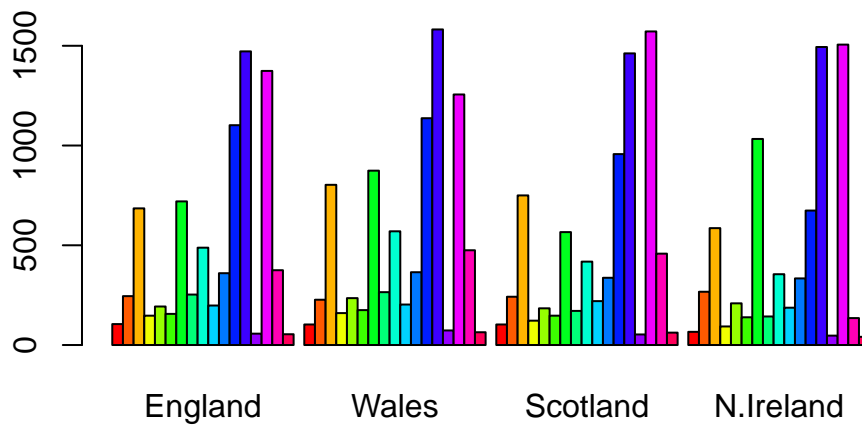
Alternatively, we could have added `row.names=1` to our initial read of the csv to do this.

Q2. Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

A2. I prefer using `row.names=1` to make changes in the original read rather than having to write out additional code to change it. Moreover, the former approach will continue to remove a column every time it is ran due to `x<-x[,-1]` will return the dataframe without the first column each time since we are not setting it to a different variable.

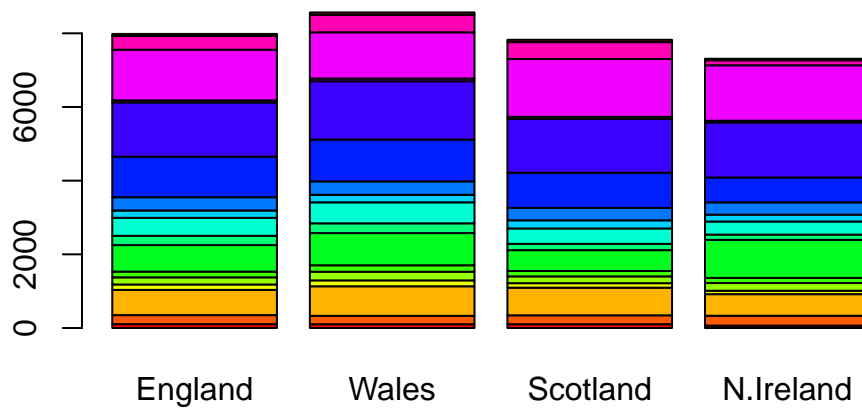
Numbers alone doesn't tell us much here. Neither do regular plots.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?


```
## Changing beside to = False will not separate each row side-by-side.
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



A3. Changing `beside=T` to `beside=F` will produce the graph above.

Q5. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
## The code in question  
pairs(x, col=rainbow(10), pch=16)
```



A5. It seems there are 4 figures that appear after each country (bar N. Ireland), but otherwise, it is quite confusing because what do those 4 figures represent. `pairs(x)` creates the scatterplot where each variable in `x` is plotted with one of the other variables for every variable. `col=rainbow(10)` sets the color to 10 different colors of the rainbow for visualization. `pch=16` uses solid circles for points. If a given point lies on the diagonal of a given plot, the correlation/relationship between the two variables of that given plot may be very strong.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

A6. The main differences are that N. Ireland eats way more `fresh_potatoes` and consumes less `alcoholic_drinks`, `cheese`, `fresh_fruit`, and `other_meat`.

PCA to visualize PCA will help us better visualize. Using `prcomp()`, we will need to transpose using `t()` because its default assumes observations is rows and variables is columns.

```
# Use the prcomp() PCA function
pca <- prcomp( t(x) )
summary(pca)
```


Importance of components:

PC1 PC2 PC3 PC4

| | | | | |
|------------------------|----------|----------|----------|-----------|
| Standard deviation | 324.1502 | 212.7478 | 73.87622 | 2.921e-14 |
| Proportion of Variance | 0.6744 | 0.2905 | 0.03503 | 0.000e+00 |
| Cumulative Proportion | 0.6744 | 0.9650 | 1.00000 | 1.000e+00 |


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
rybg<- c("orange", "red", "blue", "darkgreen")
text(pca$x[,1], pca$x[,2], colnames(x), col=rybg)
```



Below, we will calculate how much variation in the data each PC accounts for using:

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```


Which returned the proportion of variance from earlier.

```
z <- summary(pca)
z$importance
```

| | PC1 | PC2 | PC3 | PC4 |
|------------------------|-----------|-----------|----------|--------------|
| Standard deviation | 324.15019 | 212.74780 | 73.87622 | 2.921348e-14 |
| Proportion of Variance | 0.67444 | 0.29052 | 0.03503 | 0.000000e+00 |
| Cumulative Proportion | 0.67444 | 0.96497 | 1.00000 | 1.000000e+00 |

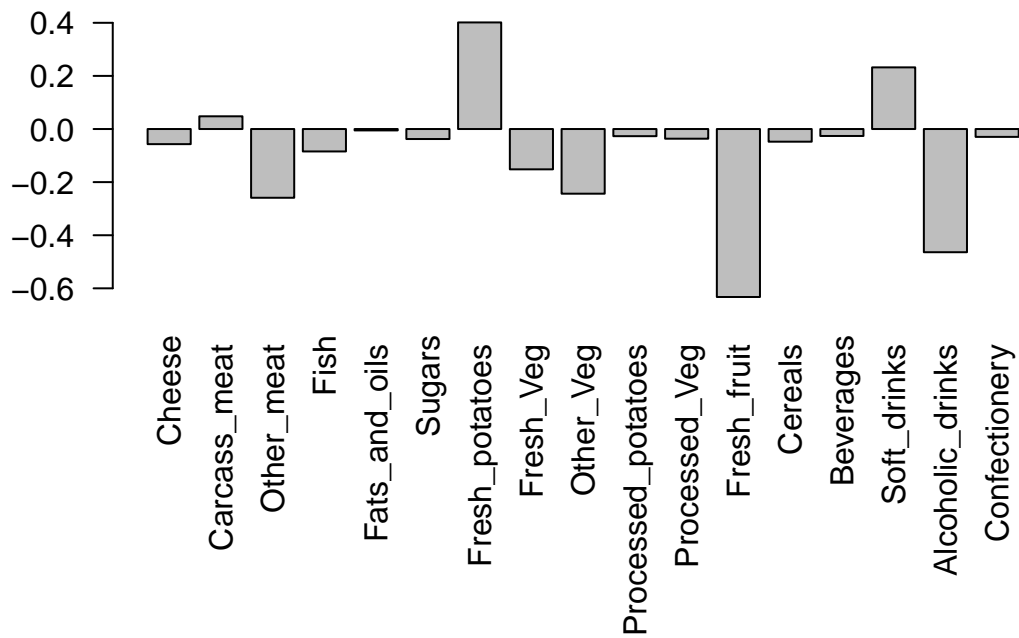
The information can be summarized in a plot of eigenvalues (variances).

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```




Moving on to variable loadings, which considers the influence of each of the original variables on the PC.

```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```

Q9. Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```




A9. The two groups that feature prominently are fresh_potatoes and soft_drinks negatively and positively pushing, respectively.

Using ggplot for the figures!

```
library(ggplot2)

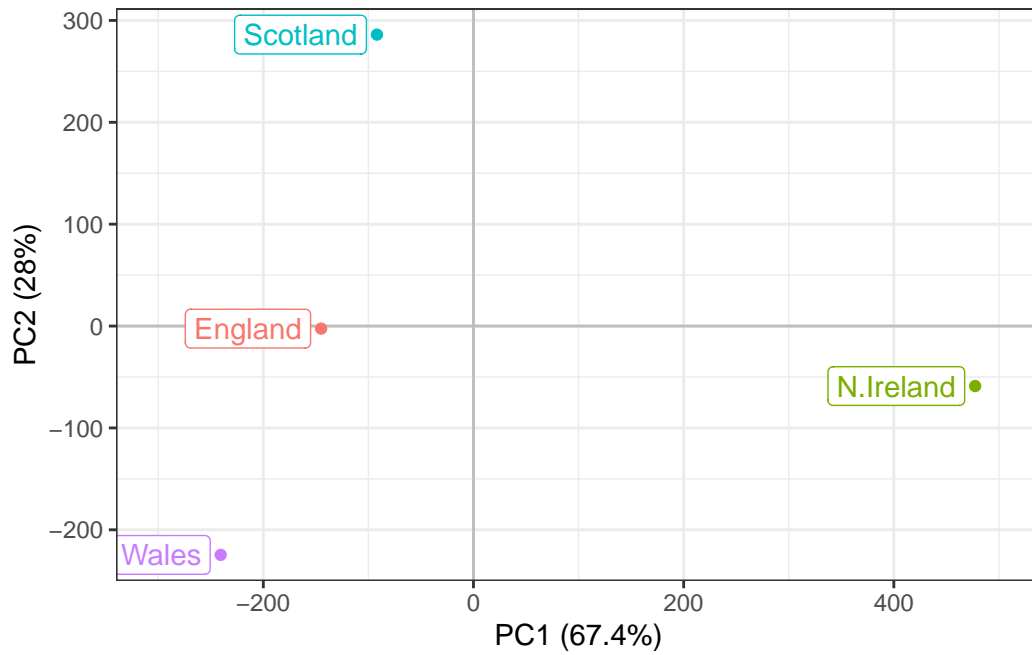
df <- as.data.frame(pca$x)
df_lab <- tibble::rownames_to_column(df, "Country")

# Our first basic plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country) +
  geom_point()
```



Adding more:

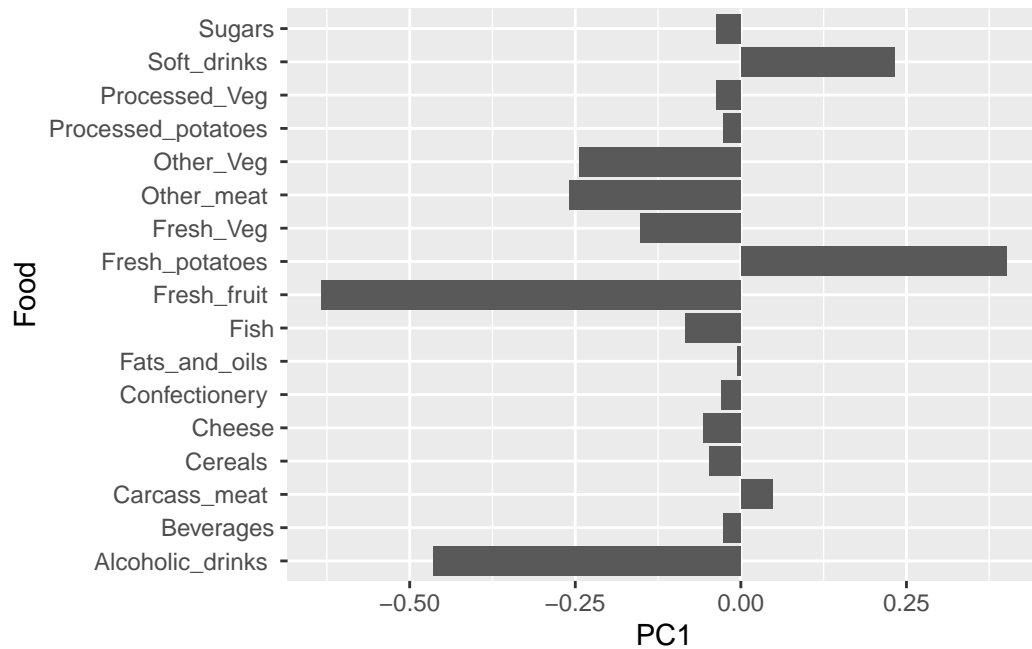
```
ggplot(df_lab) +  
  aes(PC1, PC2, col=Country, label=Country) +  
  geom_hline(yintercept = 0, col="gray") +  
  geom_vline(xintercept = 0, col="gray") +  
  geom_point(show.legend = FALSE) +  
  geom_label(hjust=1, nudge_x = -10, show.legend = FALSE) +  
  expand_limits(x = c(-300,500)) +  
  xlab("PC1 (67.4%)") +  
  ylab("PC2 (28%)") +  
  theme_bw()
```



Let's do the same with the loadings figure!


```
ld <- as.data.frame(pca$rotation)
ld_lab <- tibble::rownames_to_column(ld, "Food")

ggplot(ld_lab) +
  aes(PC1, Food) +
  geom_col()
```



Adding more:

```
ggplot(ld_lab) +
  aes(PC1, reorder(Food, PC1), bg=PC1) +
  geom_col() +
  xlab("PC1 Loadings/Contributions") +
  ylab("Food Group") +
  scale_fill_gradient2(low="purple", mid="gray", high="darkgreen", guide=NULL) +
  theme_bw()
```



PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

| | wt1 | wt2 | wt3 | wt4 | wt5 | ko1 | ko2 | ko3 | ko4 | ko5 |
|-------|------|-----|------|------|-----|-----|-----|-----|-----|-----|
| gene1 | 439 | 458 | 408 | 429 | 420 | 90 | 88 | 86 | 90 | 93 |
| gene2 | 219 | 200 | 204 | 210 | 187 | 427 | 423 | 434 | 433 | 426 |
| gene3 | 1006 | 989 | 1030 | 1017 | 973 | 252 | 237 | 238 | 226 | 210 |
| gene4 | 783 | 792 | 829 | 856 | 760 | 849 | 856 | 835 | 885 | 894 |
| gene5 | 181 | 249 | 204 | 244 | 225 | 277 | 305 | 272 | 270 | 279 |
| gene6 | 460 | 502 | 491 | 491 | 493 | 612 | 594 | 577 | 618 | 638 |

Q10: How many genes and samples are in this data set?

```
dim(rna.data)
```


```
[1] 100 10
```

A10. If genes are rows and samples are columns, there are 100 genes and 10 samples.

To make the plot:

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

## Simple un polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



How much variation?

```
summary(pca)
```

Importance of components:


| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 |
|------------------------|--------|--------|---------|---------|---------|---------|---------|
| Standard deviation | 9.6237 | 1.5198 | 1.05787 | 1.05203 | 0.88062 | 0.82545 | 0.80111 |
| Proportion of Variance | 0.9262 | 0.0231 | 0.01119 | 0.01107 | 0.00775 | 0.00681 | 0.00642 |
| Cumulative Proportion | 0.9262 | 0.9493 | 0.96045 | 0.97152 | 0.97928 | 0.98609 | 0.99251 |

| | PC8 | PC9 | PC10 |
|------------------------|---------|---------|-----------|
| Standard deviation | 0.62065 | 0.60342 | 3.345e-15 |
| Proportion of Variance | 0.00385 | 0.00364 | 0.000e+00 |
| Cumulative Proportion | 0.99636 | 1.00000 | 1.000e+00 |

PC1 is the dimension of interest (92.6%), accounting for the most variation.

To quickly plot:

```
plot(pca, main="Quick scree plot")
```



Variation!

```
## Variance captured per PC
pca.var <- pca$sdev^2


## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

Again, returning what we saw in the summary.

```
barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```


Scree Plot

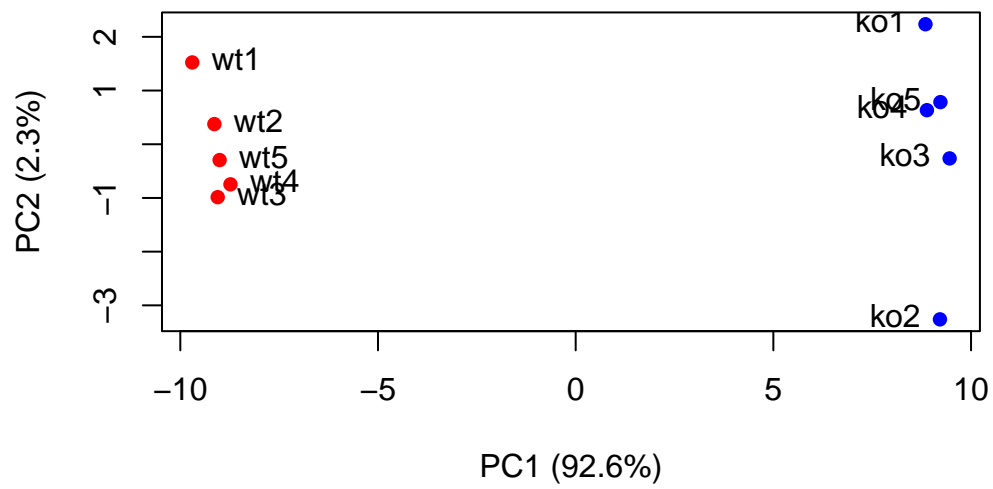


Spice it up a bit:

```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca$x[,1], pca$x[,2], col=colvec, pch=16,
     xlab=paste0("PC1 (", pca.var.per[1], "%)"),
     ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca$x[,1], pca$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```

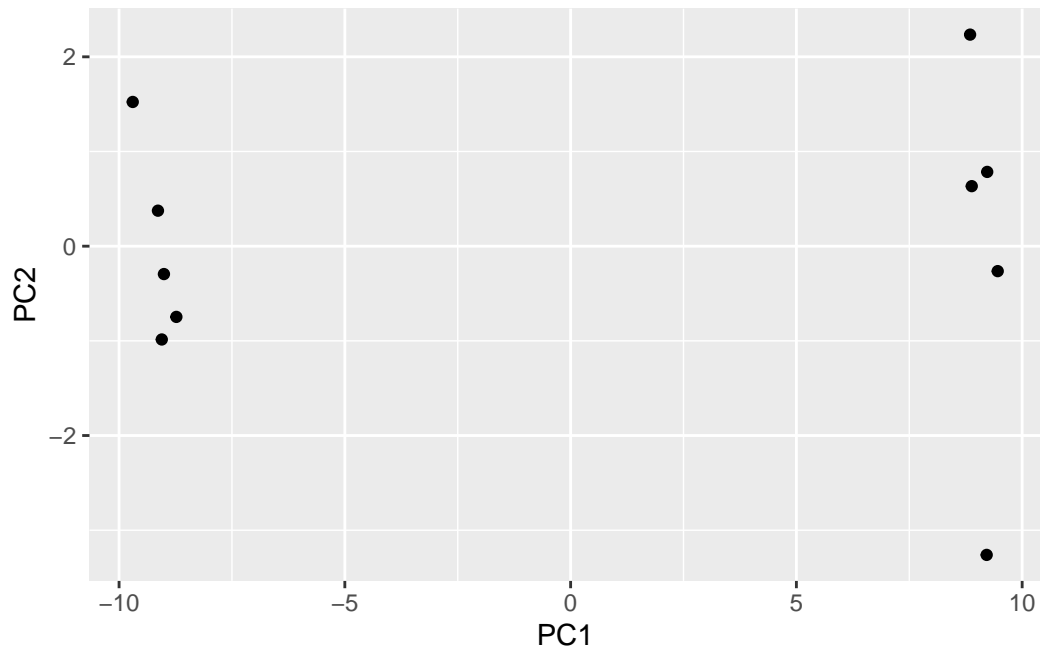


Using ggplot:

```
library(ggplot2)

df <- as.data.frame(pca$x)


# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



To add the information from our initial dataframe

```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE)
p
```




Finally

```
p + labs(title="PCA of RNASeq Data",
  subtitle = "PC1 clealy seperates wild-type from knock-out samples",
  x=paste0("PC1 (", pca.var.per[1], "%)"),
  y=paste0("PC2 (", pca.var.per[2], "%)"),
  caption="Class example data") +
  theme_bw()
```

PCA of RNASeq Data

PC1 clearly separates wild-type from knock-out samples



Class example data

What a beaut.